

IDS 572 ASSIGNMENT 2

Prediction and Investment Decisions - Lending Club

Harika Kalyani Lakhinena - hlakhi2@uic.edu - UIN: 676218205

Roshan Dhanasiri - rdhana5@uic.edu - UIN: 676134844

Sathwik Reddy Maddi - smaddi9@uic.edu - UIN: 658204431

1. (a) Develop boosted tree models (using either gbm or xgBoost) to predict loan_status. Experiment with different parameters using a grid of parameter values. Use cross-validation. Explain the rationale for your experimentation. How does performance vary with parameters, and which parameter setting you use for the 'best' model?

Model performance should be evaluated through use of same set of criteria as for the earlier models - confusion matrix based, ROC analyses and AUC, cost-based performance.



Ans. We're predicting loan_status using GBM.

GBM using Over-sampled data:

Variable Importance:

```
A gradient boosted model with bernoulli loss function.
1000 iterations were performed.
The best test-set iteration was 1000.
There were 42 predictors of which 40 had non-zero influence.
> summary(gbm_paramTune_os)
```

	var	rel.inf
sub_grade	sub_grade	29.37019887
addr_state	addr_state	26.87944626
int_rate	int_rate	14.49181093
acc_open_past_24mths	acc_open_past_24mths	2.57238301
dti	dti	2.10060938
avg_cur_bal	avg_cur_bal	1.96298874
tot_hi_cred_lim	tot_hi_cred_lim	1.88380652
purpose	purpose	1.43977922
mo_sin_old_il_acct	mo_sin_old_il_acct	1.36983110
mo_sin_old_rev_tl_op	mo_sin_old_rev_tl_op	1.36323598
loan_amnt	loan_amnt	1.32584361
total_rev_hi_lim	total_rev_hi_lim	1.18867531
mths_since_recent_bc	mths_since_recent_bc	1.15180751
bc_util	bc_util	1.10525669
revol_bal	revol_bal	1.09663839
annual_inc	annual_inc	1.02533797
bc_open_to_buy	bc_open_to_buy	1.01056551
total_bc_limit	total_bc_limit	0.79145888
total_bal_ex_mort	total_bal_ex_mort	0.69234696
pct_tl_nvr_dlq	pct_tl_nvr_dlq	0.64339673
tot_cur_bal	tot_cur_bal	0.63983038
total_il_high_credit_limit	total_il_high_credit_limit	0.63978395
revol_util	revol_util	0.63798531
total_acc	total_acc	0.52108200
num_actv_bc_tl	num_actv_bc_tl	0.48595368
home_ownership	home_ownership	0.41830992
num_rev_accts	num_rev_accts	0.40732341
num_bc_tl	num_bc_tl	0.40030714
num_rev_tl_bal_gt_0	num_rev_tl_bal_gt_0	0.33801129
num_actv_rev_tl	num_actv_rev_tl	0.33110438
num_op_rev_tl	num_op_rev_tl	0.28029623
mo_sin_rcnt_rev_tl_op	mo_sin_rcnt_rev_tl_op	0.26789317



Accuracy with Threshold (0.5) = 0.8125742
Area Under Curve = 0.6794854



GBM using Under-sampled data:

Variable Importance:

```
A gradient boosted model with bernoulli loss function.
1000 iterations were performed.
The best test-set iteration was 87.
There were 36 predictors of which 31 had non-zero influence.
> summary(gbm_paramTune_us)
```

	var	rel.inf
sub_grade	sub_grade	31.91082033
addr_state	addr_state	31.81100925
int_rate	int_rate	7.72148760
dti	dti	2.30672781
loan_amnt	loan_amnt	2.09651176
mo_sin_old_rev_tl_op	mo_sin_old_rev_tl_op	1.92037631
acc_open_past_24mths	acc_open_past_24mths	1.82493941
total_bc_limit	total_bc_limit	1.51468447
total_bal_ex_mort	total_bal_ex_mort	1.44928792
revol_bal	revol_bal	1.43803822
avg_cur_bal	avg_cur_bal	1.40294668
total_rev_hi_lim	total_rev_hi_lim	1.37367196
annual_inc	annual_inc	1.35306004
tot_cur_bal	tot_cur_bal	1.21878019
tot_hi_cred_lim	tot_hi_cred_lim	1.17070015
total_il_high_credit_limit	total_il_high_credit_limit	1.04021441
purpose	purpose	1.00434741
pct_tl_nvr_dlq	pct_tl_nvr_dlq	0.73163135
mo_sin_rcnt_rev_tl_op	mo_sin_rcnt_rev_tl_op	0.72264838
total_acc	total_acc	0.70412777
num_rev_accts	num_rev_accts	0.58682243
num_bc_tl	num_bc_tl	0.51648829
num_tl_op_past_12m	num_tl_op_past_12m	0.49943079
mort_acc	mort_acc	0.45875864
mo_sin_rcnt_tl	mo_sin_rcnt_tl	0.43478643
num_rev_tl_bal_gt_0	num_rev_tl_bal_gt_0	0.40627906
num_bc_sats	num_bc_sats	0.39660161
open_acc	open_acc	0.39582334
num_actv_bc_tl	num_actv_bc_tl	0.33811967
num_actv_rev_tl	num_actv_rev_tl	0.31001555
num_il_tl	num_il_tl	0.26820116
home_ownership	home_ownership	0.23934229
num_op_rev_tl	num_op_rev_tl	0.23447582

Accuracy VS Cutoff plot:



Accuracy with Threshold (0.5): 0.8321157


Area Under Curve = 0.7007469

Variable Importance:

```
A gradient boosted model with bernoulli loss function.
1000 iterations were performed.
The best test-set iteration was 1000.
There were 42 predictors of which 40 had non-zero influence.
> summary(gbm_paramTune)
```


	var	rel.inf
sub_grade	sub_grade	33.23943743
addr_state	addr_state	28.55039410
int_rate	int_rate	8.04243661
revol_util	revol_util	3.28600238
acc_open_past_24mths	acc_open_past_24mths	2.46693936
dti	dti	2.20867940
avg_cur_bal	avg_cur_bal	1.87844915
tot_hi_cred_lim	tot_hi_cred_lim	1.78856025
mo_sin_old_il_acct	mo_sin_old_il_acct	1.45013325
bc_util	bc_util	1.37924574
annual_inc	annual_inc	1.20493254
mths_since_recent_bc	mths_since_recent_bc	1.13849389
loan_amnt	loan_amnt	1.10484985
total_rev_hi_lim	total_rev_hi_lim	1.00333749
bc_open_to_buy	bc_open_to_buy	0.95439609
mo_sin_old_rev_tl_op	mo_sin_old_rev_tl_op	0.90363841
purpose	purpose	0.81916582
home_ownership	home_ownership	0.80627995
revol_bal	revol_bal	0.70868291
total_bal_ex_mort	total_bal_ex_mort	0.62724410
total_bc_limit	total_bc_limit	0.55805037
total_il_high_credit_limit	total_il_high_credit_limit	0.54666759
num_actv_rev_tl	num_actv_rev_tl	0.46559070
mo_sin_rcnt_tl	mo_sin_rcnt_tl	0.45332756
pct_tl_nvr_dlq	pct_tl_nvr_dlq	0.44066133
num_bc_tl	num_bc_tl	0.41815570
num_actv_bc_tl	num_actv_bc_tl	0.40651696
num_il_tl	num_il_tl	0.38056084
num_rev_tl_bal_gt_0	num_rev_tl_bal_gt_0	0.36063540
num_rev_accts	num_rev_accts	0.32503092
mort_acc	mort_acc	0.31442553
tot_cur_bal	tot_cur_bal	0.31127919
mo_sin_rcnt_rev_tl_op	mo_sin_rcnt_rev_tl_op	0.26320288
total_acc	total_acc	0.23043274

Accuracy vs Cutoff Plot:



Accuracy with Threshold (0.5) = 0. 8125742

ROC Curve:




Area under curve = 0.6751129


Provide a table with comparative evaluation of all the best models from each method (decision trees and random forests from the earlier assignment and boosted trees); show their ROC curves in a combined plot. Also provide profit-curves and 'best' profit' and associated cutoff values. At the respective best cutoff levels, what are the accuracy values for the different models?

Evaluation of Decision Trees

ROC Curve:




LIFT Curve:



AUC is 0.5440


Evaluation of Random Forests

ROC Curve:




AUC: 0.9470


Lift Curve:



Accuracy VS Cutoff plot:



Accuracy with Threshold (0.5): 0.8321157




Area Under Curve = 0.7007469

2. (a) Develop linear (glm) models to predict loan_status. Experiment with different parameter values and identify which gives 'best' performance. Use cross-validation. Describe how you determine 'best' performance. How do you handle variable selection?


Experiment with Ridge and Lasso, and show how you vary these parameters, and what performance is observed.

Ans. The best model is determined by the model which has the highest accuracy and area under the curve. In GLM the model with best performance is GLM using Ridge function.

GLM:




Area Under Curve = 0.6955761





Accuracy= 0.850937488753734

Plot of GLMNET:




GLM using Ridge:





Accuracy = 0.851117429013567

Area under curve: 0.6968744

Plot of GLMNET using Ridge:




GLMNET Using LASSO:



Accuracy = 0.851009464857667
AUC = 0.6938379

Plot of GLMNET using LASSO:




Variable selection can be handled by penalizing the magnitude of coefficients of features along with minimizing the error between predicted and actual observations. This is done by using Ridge and Lasso.

For Ridge Regression α value = 0,

```
glmDefault_cv_lasso <- cv.glmnet(data.matrix(xD), lcdfTrn$loan_status, family="binomial",
alpha = 0)
```

For Lasso regression α value = 1,

```
glmDefault_cv_lasso <- cv.glmnet(data.matrix(xD), lcdfTrn$loan_status, family="binomial", alpha
= 1)
```



(b) For the linear model, what is the loss function, and link function you use? (Write the expression for these, and briefly describe).

Ans.

RIDGE:

Loss function = OLS + alpha * summation (squared coefficient values)

loss function is sum of squared errors

L2 regularization

We can automate this task of finding the optimal lambda value using the `cv.glmnet()` function. The optimal lambda value comes out to be -2.1 and will be used to build the ridge regression model.

LASSO:

Loss function = OLS + alpha * summation (absolute values of the magnitude of the coefficients)

L1 norm -- can give sparse models

In the above function, alpha is the penalty parameter we need to select. Using an l1-norm constraint forces some weight values to zero to allow other coefficients to take non-zero values.

The first step to build a lasso model is to find the optimal lambda value using the code below. For lasso regression, the alpha value is 1. The output is the best cross-validated lambda, which comes out to be -4.9.

The loss function defined in the 2-class classification problem is known as Cross-Entropy loss function.

$$y^t \in 0, 1 \quad y^t \sim \text{Ber}(y^t; \mu^t) \quad 0 \leq \mu^t \leq 1$$

$$p(y|X, w) = \prod_{t=1}^N (\mu^t)^{y^t} (1 - \mu^t)^{(1-y^t)}$$

$$L = -\log p(y|X, w) = \sum_t -y^t \log \mu^t - (1 - y^t) \log(1 - \mu^t)$$

Cross-Entropy Loss Function

Link function = $\ln(\mu/(1-\mu))$

(c) Compare performance of models with that of random forests (from last assignment) and gradient boosted tree models.

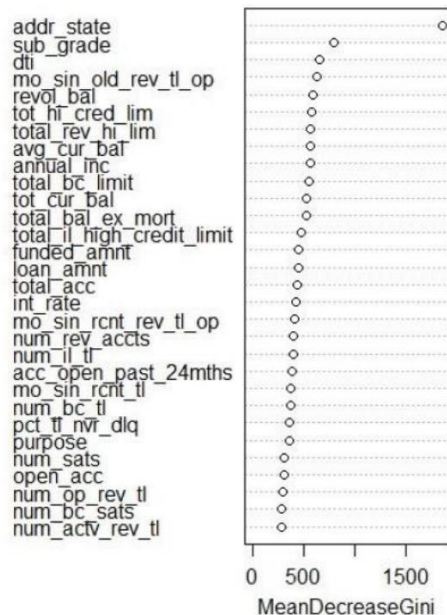
Ans. By comparing the performance of models with random forests, GBM performed better in terms of accuracy and area under curve.

(d) Examine which variables are found to be important by the best models from the different methods, and comment on similarities, difference. What do you conclude?

Ans. For all the best models with the best different methods, most of the variables which are important are the same. But, the order of their importance differed.

For Random Forest, based on mean decrease Gini:

Addr_state was most important, followed by sub grade, dti, mo_sin_old_rev_tl_op, revol_bal, tot_hi_cred_lim, total_rev_hi_lim, avg_cur_bal, annual_inc, total_bc_limit, tot_cur_bal, total_bal_ex_mort and so on upto num_actv_rev_tl.



RANDOM FOREST

var	rel.inf
sub_grade	31.91082033
addr_state	31.81100925
int_rate	7.72148760
dti	2.30672781
loan_amnt	2.09651176
mo_sin_old_rev_tl_op	1.92037631
acc_open_past_24mths	1.82493941
total_bc_limit	1.51468447
total_bal_ex_mort	1.44928792
revol_bal	1.43803822
avg_cur_bal	1.40294668
total_rev_hi_lim	1.37367196
annual_inc	1.35306004
tot_cur_bal	1.21878019
tot_hi_cred_lim	1.17070015
total_il_high_credit_limit	1.04021441
purpose	1.00434741
pct_tl_nvr_dlq	0.73163135
mo_sin_rcnt_rev_tl_op	0.72264838
total_acc	0.70412777
num_rev_accts	0.58682243
num_bc_tl	0.51648829
num_tl_op_past_12m	0.49943079
mort_acc	0.45875864
mo_sin_rcnt_tl	0.43478643
num_rev_tl_bal_gt_0	0.40627906
num_bc_sats	0.39660161
open_acc	0.39582334
num_actv_bc_tl	0.33811967
num_actv_rev_tl	0.31001555
num_il_tl	0.26820116
home_ownership	0.23934229
num_op_rev_tl	0.23447582

GBM Using Under Sampling Data

Label	Value
(Intercept)	2.9448240000
mort_acc	0.0160469700
mo_sin_rcnt_tl	0.0033564710
mo_sin_rcnt_rev_tl_op	0.0014692550
pct_tl_nvr_dlq	0.0011798440
open_acc	0.0011718050
num_sats	0.0010358740
addr_state	0.0008879191
num_rev_accts	0.0006525662
total_acc	0.0005702526
purpose	0.0005104390
mo_sin_old_rev_tl_op	0.0002690454
total_bc_limit	0.0000025714
avg_cur_bal	0.0000015161
total_rev_hi_lim	0.0000010336
annual_inc	0.0000002686
tot_hi_cred_lim	0.0000002091
tot_cur_bal	0.0000001845
total_bal_ex_mort	0.0000001284
revol_bal	0.0000000278
loan_amnt	-0.0000026857
funded_amnt	-0.0000026873
num_il_tl	-0.0001966692
num_op_rev_tl	-0.0012443990
num_bc_tl	-0.0014236930
num_bc_sats	-0.0022205840
num_actv_rev_tl	-0.0073991720
dti	-0.0080257270
num_rev_tl_bal_gt_0	-0.0090145210
num_actv_bc_tl	-0.0108421700
num_tl_op_past_12m	-0.0244462200
sub_grade	-0.0245549400
acc_open_past_24mths	-0.0260111800
int_rate	-0.0365062900
home_ownership	-0.0500243200
grade	-0.1127102000

GLM Using Ridge

Similarities and Differences:

Addr_state, sub_grade, dti, mo_sin_old_rev_tl_op, revol_bal, tot_hi_cred_lim, total_rev_hi_lim, avg_cur_bal, annual_inc, total_bc_limit, tot_cur_bal, total_bal_ex_mort and so on upto num_actv_rev_tl are important variables in Random Forest and GBM. But, In GLM mort_acc, mo_sin_rcnt_tl, mo_sin_rcnt_rev_tl_op, pct_tl_nvr_dlq, open_acc, num_sats, addr_state

In GLM, loan amount and interest rate are penalized. Whereas, those were two of the most important variables in GBM and Random Forest.

(e) In developing models above, do you find larger training samples to give better models? Do you find balancing the training data examples across classes to give better models?


Ans. The performance of the model was presented in the table below for all GLM, GLM ridge, GLM lasso, GLM elastic models. The best performance was determined based on the values of accuracy, AUC of ROC curves.

Models which were built using under sampling data performed better when compared to all the other models. So, the models with larger training samples do not always give better models. The results obtained are pertaining to the dataset we have and it varies with the objective of our model. Yes, Balancing training data across classes gives better models. Imbalanced data affects random forests more than GBM and GLMs.


3. Develop models to identify loans which provide the best returns. Explain how you define returns? Does it include Lending Club's service costs? 2 Develop glm, rf, gbm/xgb models for this. Show how you systematically experiment with different parameters to find the best models. Compare model performance – explain what performance criteria do you use, and why.

Ans. Out of all the models whose performance was studied, it was found that the GLMNET performed best with lowest error rate for training data. Below is the model wise performance for each.

Training Plot:



Test Plot:

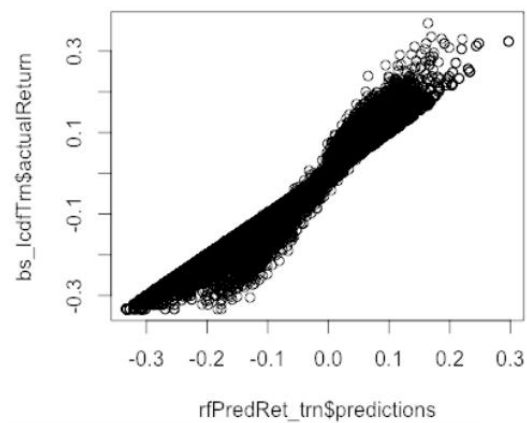


Number of Trees = 100


Training error = 0.02182912

Test Error = 0.09452704

Training plot:



Test Plot:

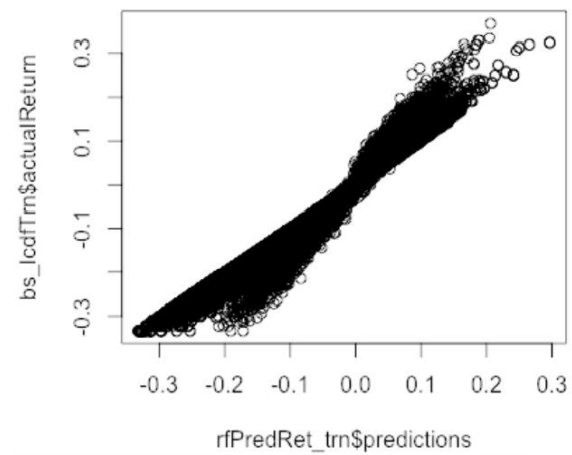


Number of trees = 200

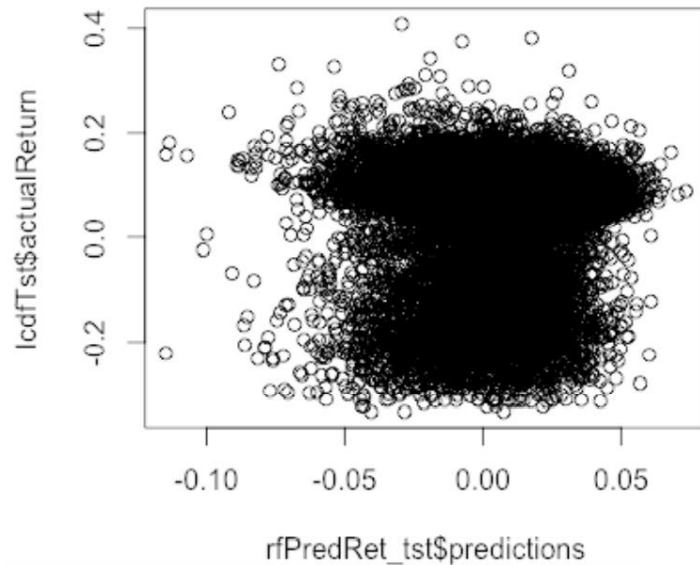
Training error = 0.021434

Test Error = 0.09426395

Training Plot:



Test Plot:



We conclude that the RM model with 200 trees to be has least test error and highest accuracy.

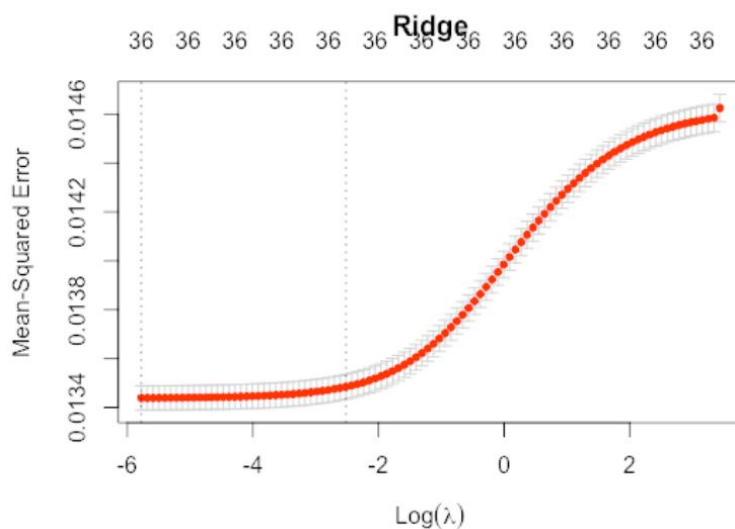
GLMNET

We tried to build generalized linear model with trying out different values of “alpha”

Starting out with alpha = 0 i.e. Ridge regression

Mean square error = 0.01202288

Plot of mean square error is:

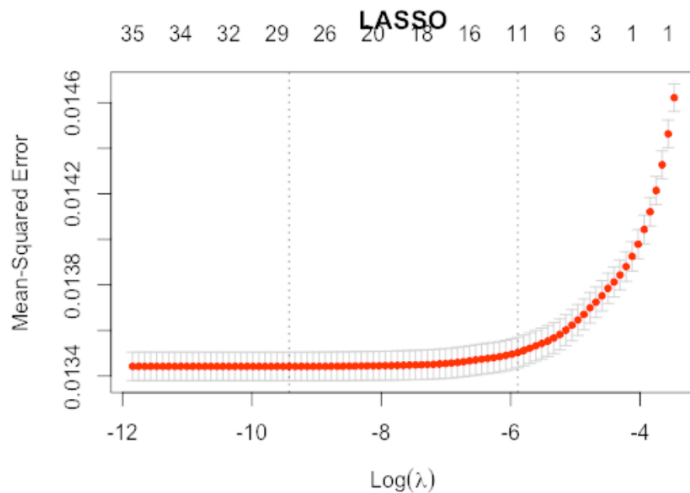


For alpha = 0.5 i.e. Elastic net

Mean square error = 0.01209347

For alpha = 1 i.e. Lasso Mean square error = 0.01209347

Plot of mean squared error : 0.01211678



So my highest accuracy model is “Ridge” with least MSE.

Gradient boosting method

We performed a grid search for hyper parameter tuning with 500 trees.



Training error for the model was :

```
error <- sqrt(mean((gbPredRet_trn- bs_1cdfTrn$actualReturn)^2))
> error
[1] 0.1137398
```

Test error for the model was:

```
error_tst <- sqrt(mean((gbPredRet_tst- 1cdfTst$actualReturn)^2))
> error_tst
[1] 0.08973837
```

Performance plot for GBM:



```
> summary(gbm_paramTune)
```

	var	rel.inf
addr_state	addr_state	30.62544037
sub_grade	sub_grade	27.95179956
annual_inc	annual_inc	5.74254655
acc_open_past_24mths	acc_open_past_24mths	2.69086877
revol_bal	revol_bal	2.61770478
total_bal_ex_mort	total_bal_ex_mort	2.27467993
bc_open_to_buy	bc_open_to_buy	2.07526714
tot_hi_cred_lim	tot_hi_cred_lim	1.80513180
total_bc_limit	total_bc_limit	1.80026566
total_rev_hi_lim	total_rev_hi_lim	1.73236672
purpose	purpose	1.70694157
mo_sin_old_rev_tl_op	mo_sin_old_rev_tl_op	1.60879852
dti	dti	1.60180613
revol_util	revol_util	1.56046142
total_il_high_credit_limit	total_il_high_credit_limit	1.49611721
avg_cur_bal	avg_cur_bal	1.19131663
mths_since_recent_bc	mths_since_recent_bc	1.17916052
bc_util	bc_util	1.12540049
total_acc	total_acc	1.09275728
tot_cur_bal	tot_cur_bal	1.07742995
int_rate	int_rate	0.88232392
num_bc_tl	num_bc_tl	0.72359706
num_il_tl	num_il_tl	0.70294777
num_rev_accts	num_rev_accts	0.49450210
mo_sin_rcnt_rev_tl_op	mo_sin_rcnt_rev_tl_op	0.46094296
num_tl_op_past_12m	num_tl_op_past_12m	0.44649566
pct_tl_nvr_dlq	pct_tl_nvr_dlq	0.44310162
mo_sin_rcnt_tl	mo_sin_rcnt_tl	0.43216962
mort_acc	mort_acc	0.37704537
open_acc	open_acc	0.33499090
num_sats	num_sats	0.30742389
num_actv_rev_tl	num_actv_rev_tl	0.28493975
num_op_rev_tl	num_op_rev_tl	0.24469882
num_actv_bc_tl	num_actv_bc_tl	0.24274424
home_ownership	home_ownership	0.23408040
num_bc_sats	num_bc_sats	0.18177685
num_rev_tl_bal_gt_0	num_rev_tl_bal_gt_0	0.16980075
grade	grade	0.08015736
num_tl_120dpd_2m	num_tl_120dpd_2m	0.00000000

4. Considering results from Questions 1 and 2 above – that is, considering the best model for predicting loan-status and that for predicting loan returns -- how would you select loans for investment? There can be multiple approaches for combining information from the two models - describe your approach and show performance. How does performance here compare with use of single models?

Ans.

Decile performances by each model

1. Decile performance in RANDOM FOREST

```
# A tibble: 10 x 14
```

	tile	count	avgpredRet	numDefaults	avgActRet	minRet	maxRet	avgTer	totA	totB	totC	totD	totE	totF
	<int>	<int>	<dbl>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<int>	<int>	<int>	<int>	<int>
1	1	2779	0.0394	221	0.0556	-0.322	0.257	2.13	738	1249	600	167	24	1
2	2	2779	0.0297	237	0.0475	-0.313	0.260	2.21	1160	1083	392	120	23	1
3	3	2779	0.0244	283	0.0449	-0.300	0.221	2.27	1162	1023	438	136	19	0
4	4	2778	0.0195	279	0.0478	-0.313	0.256	2.25	1040	990	565	152	30	1
5	5	2779	0.0145	332	0.0458	-0.333	0.381	2.28	887	976	664	210	38	4
6	6	2779	0.00925	363	0.0459	-0.312	0.283	2.28	708	1015	727	265	61	3
7	7	2778	0.00323	429	0.0463	-0.310	0.319	2.29	518	933	917	344	59	5
8	8	2779	-0.00403	544	0.0439	-0.322	0.286	2.28	327	767	1107	442	124	11
9	9	2779	-0.0136	637	0.0419	-0.322	0.288	2.32	172	570	1213	618	179	22
10	10	2778	-0.0337	802	0.0367	-0.333	0.408	2.34	61	248	1022	898	425	108

We observe that in the first 7 deciles the average predicted return is positive and in the top 4 deciles the difference between the average prediction and average actual is less. The number of defaults are not significantly differing as we move down the decile and seems like grade A and B have good investment opportunities as expected and up to some extent grade C as well

2. Decile performance with GLMNET

```
# A tibble: 10 x 14
```

	tile	count	avgpredRet	numDefaults	avgActRet	minRet	maxRet	avgTer	totA	totB	totC	totD	totE	totF
	<int>	<int>	<dbl>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<int>	<int>	<int>	<int>	<int>
1	1	2779	0.0259	95	0.0402	-0.281	0.205	2.25	2209	511	53	6	0	0
2	2	2779	0.00873	155	0.0438	-0.283	0.216	2.28	1677	942	153	6	1	0
3	3	2779	-0.000232	218	0.0448	-0.333	0.233	2.22	1189	1290	276	21	3	0
4	4	2778	-0.00763	279	0.0472	-0.309	0.232	2.25	812	1410	509	43	4	0
5	5	2779	-0.0146	310	0.0509	-0.322	0.272	2.24	480	1430	765	96	7	1
6	6	2779	-0.0211	355	0.0542	-0.333	0.308	2.21	244	1240	1096	188	11	0
7	7	2778	-0.0232	467	0.0514	-0.333	0.246	2.25	121	978	1337	311	30	1
8	8	2779	-0.0361	596	0.0458	-0.333	0.289	2.28	43	629	1494	541	68	4
9	9	2779	-0.0462	694	0.0449	-0.323	0.319	2.32	5	301	1330	929	204	10
10	10	2778	-0.0661	914	0.0378	-0.333	0.502	2.32	5	62	679	1181	706	127

If we go GLMNET for decile prediction we see that only top 2 deciles give the positive average predicted return with low number of defaults with majority chunk of grade A in top 2 deciles, so I would invest in them

3. Decile performance of Gradient boosting Method (GBM):

```
# A tibble: 10 x 14
```

	tile	count	avgpredRet	numDefaults	avgActRet	minRet	maxRet	avgTer	totA	totB	totC	totD	totE	totF
	<int>	<int>	<dbl>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<int>	<int>	<int>	<int>	<int>
1	1	2779	0.0233	170	0.0493	-0.313	0.248	2.13	1348	999	357	75	0	0
2	2	2779	0.0226	270	0.0472	-0.313	0.272	2.22	999	1139	481	158	1	0
3	3	2779	0.0212	280	0.0473	-0.333	0.246	2.20	1023	1016	522	193	25	0
4	4	2778	0.0206	238	0.0439	-0.323	0.255	2.28	1127	1406	210	27	7	1
5	5	2779	0.0201	281	0.0431	-0.323	0.233	2.30	999	1376	314	77	13	0
6	6	2779	0.0196	358	0.0425	-0.314	0.287	2.33	759	1352	523	124	20	0
7	7	2778	0.0188	458	0.0437	-0.323	0.335	2.30	405	929	1073	325	41	4
8	8	2779	0.0175	563	0.0439	-0.322	0.305	2.27	210	554	1408	540	58	7
9	9	2779	0.0152	706	0.0454	-0.333	0.375	2.29	0	0	1652	890	218	16
10	10	2778	0.0109	863	0.0408	-0.333	0.502	2.35	0	0	1178	883	588	111

The most stable and proportionate decile performance is given by GBM. Other than obvious investing in grade A and grade B we can also go with grade C in the 6th decile. As always we observe that the maximum returns are for the bottom decile because of high risk as they are more prone to default and hence high risk with high returns.

So, I would choose to go with GBM decile performance as it gives me more proportionate observations to choose my investment from and each decile has positive average prediction.

5. As seen in data summaries and your work in the first assignment, higher grade loans are less likely to default, but also carry lower interest rates; many lower grad loans are fully paid, and these can yield higher returns. One approach may be to focus on lower grade loans (C and below) and try to identify those which are likely to be paid off. Develop models from the data on lower grade loans and check if this can provide an effective investment approach – for this, you can use one of the methods (glm, rf, or gbm/xgb) which you find to give superior performance from earlier questions. Can this provide a useful approach for investment? Compare performance with that in Question 4.

Ans.

Grade A and Grade B definitely are safe investments with low risk.

We trained our model on the lower loan grades, i.e. all the loans graded lower than B.

I checked the decile performance from a **Random Forest** trained model and found the results shown below:

```
# A tibble: 10 x 14
```

	tile	count	avgpredRet	numDefaults	avgActRet	minRet	maxRet	avgTer	totA	totB	totC	totD	totE	totF
	<int>	<int>	<dbl>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<int>	<int>	<int>	<int>	<int>
1	1	6484	0.0970	159	0.122	0.0744	0.331	1.43	45	914	2633	2043	716	101
2	2	6484	0.0653	273	0.0855	0.0563	0.196	2.04	343	2168	2896	858	195	23
3	3	6484	0.0522	346	0.0696	0.0435	0.154	2.24	1073	3106	1766	454	80	4
4	4	6483	0.0423	411	0.0566	0.0329	0.132	2.46	2207	2861	1112	245	56	2
5	5	6484	0.0329	612	0.0441	0.0164	0.123	2.67	3919	1696	636	168	63	2
6	6	6484	0.00234	4763	0.00448	-0.0914	0.115	2.87	1689	1679	1851	893	311	50
7	7	6483	-0.0008	6483	-0.0008	-0.206	-0.0252	3	812	1872	2201	1135	379	71
8	8	6484	-0.120	6484	-0.133	-0.288	-0.0989	3	730	1580	2385	1225	459	100
9	9	6484	-0.174	6484	-0.189	-0.313	-0.148	3	486	1472	2416	1467	541	91
10	10	6483	-0.237	6483	-0.246	-0.333	-0.201	3	364	1353	2306	1650	608	162

We see that the bottom 3 deciles are almost all defaulters and we observe that the average prediction is more than the average of actual. Most C and D grade are loans are pushed to the top decile which was expected and the major relative proportion of grade F are in top decile as well so maybe we can focus on these people more for high reward.

But when I checked the decile performance using GLMNET I could not see major difference as all the top loans consisted of Grade A and Grade B.

```
# A tibble: 10 x 14
```

	tile	count	avgpredRet	numDefaults	avgActRet	minRet	maxRet	avgTer	totA	totB	totC	totD	totE	totF
	<int>	<int>	<dbl>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<int>	<int>	<int>	<int>	<int>
1	1	6484	0.0205	1238	0.0175	-0.323	0.190	2.37	4730	1520	223	11	0	0
2	2	6484	0.00331	2046	0.00343	-0.313	0.208	2.45	3091	2820	525	48	0	0
3	3	6484	-0.00627	2349	-0.00230	-0.333	0.249	2.48	1972	3319	1067	126	0	0
4	4	6483	-0.0141	2816	-0.0112	-0.322	0.238	2.51	1012	3357	1864	230	20	0
5	5	6484	-0.0212	3173	-0.0179	-0.313	0.234	2.55	495	2779	2706	479	25	0
6	6	6484	-0.0282	3496	-0.0251	-0.333	0.277	2.60	219	2135	3317	759	49	5
7	7	6483	-0.0352	3882	-0.0394	-0.333	0.318	2.63	115	1463	3450	1298	149	8
8	8	6484	-0.0430	4120	-0.0462	-0.333	0.270	2.66	27	922	3296	1860	368	11
9	9	6484	-0.0528	4466	-0.0576	-0.333	0.317	2.71	1	312	2676	2607	845	40
10	10	6483	-0.0724	4912	-0.0702	-0.333	0.331	2.75	6	74	1078	2720	1952	542

6. Considering all your results, which approach(s) would you recommend for investing in LC loans? Explain your rationale.

Ans.

From the standpoint of an investor, we can see that they are limited by both money and time. We should employ alternative models that suit their risk appetite based on their risk aversion.

As a result of the above analysis,

We've noticed that GBM appears to perform better, although it demands a lot of computing power.

Based on their risk-return appetite, there are three sorts of investors: Low, Medium, and High.

In terms of Low,

The GBM loan status model will be used to recommend outcomes. Because it provides us with the best loans possible.

There's a good chance you'll get paid. Despite the lower returns on these loans, the investors in this group are confident.

Higher returns aren't an issue.

In terms of Medium,

We will suggest the combination model that we developed in response to question 5. The default rate increases somewhat, but the returns increase dramatically. This combination will benefit this type of investment the most.

of potential investors

In terms of High,

We shall suggest a model based on lower-quality loans. We've noticed that the loans in the top tile of this one yields the best results, when compared to the other ways. This category's investors prioritize high returns, therefore this model is ideal for them.